

Facial Expression Recognition via Logistic Regression and K-Nearest Neighbors

1. Introduction

In this project, we'll classify facial expressions into five categories using two machine learning algorithms:

- Logistic Regression
- K-Nearest Neighbors (KNN) Classifier

We will use a publicly available dataset, preprocess the images, extract relevant features, and evaluate the models based on various classification metrics.

2. Dataset Description

General Information

- DataSet Name: Facial Expression Recognition Challenge (FER 2013)
- Source: Kaggle Competition - Challenges in Representation Learning: Facial Expression Recognition Challenge

Data Summary

- Total Number of Samples in Dataset: 28,709 images
- Classes and Their Labels:

Label	Emotion
0	Angry
1	Disgust
2	Fear
3	Happy
4	Sad
5	Surprise
6	Neutral

Selected Classes for This Project

We removed the classes labeled as 0 for Angry and 1 for Disgust to fulfill the project requirement of using only up to 5 classes.

- Classes Used:

Label	Emotion
-------	---------

2	Fear
---	------

3	Happy
---	-------

4	Sad
---	-----

5	Surprise
---	----------

6	Neutral
---	---------

Image Specifications

- Image Size: Original images are 48x48 pixels in grayscale.
- Processed Image Size: Resized to 64x64 pixels for feature extraction.

Data Split

- Total Processed Images: 22,063
- Training Set Size: 17,650 images
- Testing Set Size: 4,413 images
- Validation Set: Not used in this project.

3. Data Preprocessing

We performed several preprocessing steps to prepare the data for modeling.

Filtering Classes

We removed the classes with labels 0 and 1.

```
# Exclude classes 0 and 1
df = df[df['emotion']!= 0]
df = df[df['emotion']!= 1]

# Check unique emotion labels
df["emotion"].unique()

Output:
array([2, 4, 6, 3, 5])
```

No missing values were found in the dataset.

4. Exploratory Data Analysis

Visualizing Sample Images

We visualized a subset of images to get an understanding of the data.

```
emotion_dict = {2: 'Fear', 3: 'Happy', 4: 'Sad', 5: 'Surprise', 6: 'Neutral'}  
fig = plt.figure(figsize=(8, 8))  
  
rows = 3  
  
columns = 3  
  
pixels = np.fromstring(subset_df['pixels'].iloc[i], sep=' ')  
  
plt.show()
```

5. Feature Extraction

Preparing Data for Feature Extraction

```
pixels = df['pixels']  
  
labels = df['emotion']
```

5.2 HOG Feature Extraction

We used Histogram of Oriented Gradients (HOG) to extract features from the images.

```
# HOG Feature Extraction Function  
  
def extract_hog_features(image):  
    pixels_per_cell=(8, 8),  
    cells_per_block=(2, 2),  
    orientations=9,  
    visualize=False)  
  
    return features
```

5.3 Processing Images and Extracting Features

To manage computational resources and balance the dataset, we limited the number of images per emotion to 5,000.

1Processed 22063 images.

6. Model Implementation

6.1 Splitting the Data

6.2 Logistic Regression

6.2.1 Training the Model

Output:

Logistic Regression - Train accuracy: 0.6107648725212464

Logistic Regression - Test accuracy: 0.48810333106730114

6.3 K-Nearest Neighbors

6.3.1 Training the Model

6.3.2 Predictions

Output:

KNN - Train accuracy: 0.5741643059490085

KNN - Test accuracy: 0.489689553591661

7. Model Evaluation

7.1 Logistic Regression Evaluation

7.1.1 Confusion Matrix

7.1.2 Classification Report

Output (Training Data):

Classification Report - Logistic Regression (Training Data):

		precision	recall	f1-score	support
2					
3					
4	2	0.56	0.46	0.50	3277
5	3	0.66	0.74	0.70	4001
6	4	0.54	0.55	0.54	3863
7	5	0.70	0.70	0.70	2534
8	6	0.59	0.62	0.60	3975
9					
10	accuracy			0.61	17650
11	macro avg	0.61	0.61	0.61	17650
12	weighted avg	0.61	0.61	0.61	17650

Output (Test Data):

Classification Report - Logistic Regression (Test Data):

2 precision recall f1-score support

3

4 2 0.37 0.31 0.34 820

5 3 0.58 0.64 0.61 999

6 4 0.41 0.42 0.41 967

7 5 0.59 0.57 0.58 637

8 6 0.48 0.50 0.49 990

9

10 accuracy 0.49 4413

11 macro avg 0.49 0.49 0.49 4413

12 weighted avg 0.48 0.49 0.49 4413

7.1.3 ROC and AUC

7.2 K-Nearest Neighbors Evaluation

7.2.1 Confusion Matrix

7.2.2 Classification Report

Output (Training Data):

1 Classification Report - KNN (Training Data):

2 precision recall f1-score support

3

4 2 0.57 0.36 0.44 3277

5 3 0.56 0.84 0.67 4001

6 4 0.60 0.40 0.48 3863

7 5 0.67 0.63 0.65 2534

8 6 0.53 0.62 0.57 3975

9

10 accuracy 0.57 17650

11 macro avg 0.59 0.57 0.56 17650

12 weighted avg 0.58 0.57 0.56 17650

Test Data:

Classification Report - KNN (Test Data):

		precision	recall	f1-score	support
2					
3					
4	2	0.44	0.27	0.34	820
5	3	0.52	0.80	0.63	999
6	4	0.49	0.29	0.36	967
7	5	0.57	0.55	0.56	637
8	6	0.43	0.51	0.47	990
9					
10	accuracy			0.49	4413
11	macro avg	0.49	0.49	0.47	4413
12	weighted avg	0.49	0.49	0.47	4413

7.2.3 ROC and AUC

8. Results and Discussion

8.1 Comparison of Models

Metric	Logistic Regression	K-Nearest Neighbors
--------	---------------------	---------------------

Training Accuracy	0.6108	0.5742
-------------------	--------	--------

Testing Accuracy	0.4881	0.4897
------------------	--------	--------

8.2 Observations

- Accuracy:
- Both models have testing accuracy of ~49%.
- Training accuracy is higher than testing accuracy, which means there is a possibility of overfitting.
- Confusion Matrices:
- The confusion matrices show that some classes are better predicted than others.
- Class 3 (Happy) has higher recall and precision across both models.
- Classification Reports:
- Precision and recall vary across classes.
- Class 3 consistently shows better performance.
- ROC Curves and AUC:
- ROC curves provide a graphical representation of the models' performance across different thresholds.

- Because of the multiclass nature and possible warnings, interpretation might be complex.

8.3 Discussion

- Model Performance:
 - Logistic Regression performed slightly better on training data.
 - KNN had comparable performance on test data.
 - The models might be affected by class imbalance or limited feature representation.
- Feature Extraction:
 - HOG features capture important structural information but may not be sufficient alone.
 - Considering other feature extraction methods or deep learning approaches might improve performance.
- Class Imbalance:
 - Though we capped the number of images per class, there may still be some imbalance.
 - Data augmentation or class weighting techniques could be useful.

9. Conclusion

In this project, we developed and implemented Logistic Regression and K-Nearest Neighbors classifiers for recognizing facial expressions in grayscale images. We preprocessed the data, extracted HOG features, and evaluated the models using different metrics for classification.

Key Observations:

- Both models had an accuracy of approximately 49% for the test data.
- Class 3, or Happy, was most predictable while the other classes showed low performance.
- There is much scope for improving feature extraction and the choice of models.

Future Scope:

- Deep learning models like Convolutional Neural Networks (CNNs) could be explored since these perform well in image classification problems.
- Data augmentation will be done to further diversify the dataset.
- Perform hyperparameter tuning. More feature extraction techniques may be tried.